

Individual Assignment 2

Weilong Chen + weilong

Problem 1

A

The provided code is an example of a multi-tasking system using FreeRTOS, where three tasks (Task 1, Task 2, and Task 3) are created and scheduled to run concurrently. The tasks are synchronized using two mutex semaphores, *resource_a* and *resource_b*.

A deadlock can occur when multiple tasks acquire resources in a way that creates a circular dependency, preventing any task from proceeding. Let's analyze the code to identify potential deadlock scenarios.

In the given code, there is a possibility of a deadlock if the following conditions are met:

Task A acquires *resource_a* and then attempts to acquire *resource_b*. Task B acquires *resource_b* and then attempts to acquire *resource_a*. If both Task A and Task B reach these points simultaneously, a circular dependency is created:

Task A holds *resource_a* and waits for *resource_b*. Task B holds *resource_b* and waits for *resource_a*. Both tasks will be stuck indefinitely, unable to proceed, resulting in a deadlock. For figure 1, if one task is at state p1 and another at state p4. The deadlock will happen as below:

Task 1: $p0 \rightarrow p1$ take r_A

Task 1: $p1 \rightarrow p2$ take r_B

Task 1: $p2 \rightarrow p3$

Task 1: $p3 \rightarrow p4$ give r_A

Task 2: $p0 \rightarrow p1$ take r_A

Deadlock

B

In a deadlock situation, one or more processes are blocked while requesting resources that are held by other processes, creating a chain of dependencies. This chain forms a cycle because each blocked process is waiting for a resource held by another blocked process, resulting in a perpetual loop of blocked states. The resource-allocation graph for the deadlock situation is shown in Figure 2.




Figure 1: Example of a Petri net built for a part of the BIND software




Figure 2

C

A resource C is added to avoid deadlock as shown in Figure 3. The program is updated on github.




Figure 3

Problem 2

B

The Liu-Layland criteria:

$U \equiv \sum_{i=1}^N \frac{C_i}{T_i} \leq N (2^{1/N} - 1)$ exhibit complete reliability when applied to a single task. However, when dealing with multiple tasks, the criteria can only determine if the tasks are schedulable but cannot guarantee the opposite scenario. For our case, it classifies 5 cases as SCHED_UNKNOWN.

The response time analysis is implemented in Python, which is shown below.

```

1 def calculate_response_time(task_list):
2     schedulable = "SCHED_YES"
3
4     for i in range(len(task_list)):
5         response_old = 0
6         response_new = task_list[i]["WCET"]
7
8         while response_new != response_old:
9             response_old = response_new
10            response_new = task_list[i]["WCET"]
11
12            for j in range(len(task_list)):

```

```

13         if task_list[j]["priority"] > task_list[i]["
priority"]:
14             response_new += (
15                 (response_old // task_list[j]["period"]) +
16                 1
17             ) * task_list[j]["WCET"]
18         if response_new > task_list[i]["deadline"]:
19             schedulable = "SCHED_NO"
20             break
21
22     return schedulable

```

Listing 1: response time analysis

```

1 Response time test 1: SCHED_YES
2 Response time test 2: SCHED_NO
3 Response time test 3: SCHED_NO
4 Response time test 4: SCHED_YES
5 Response time test 5: SCHED_YES

```

Listing 2: output